

Question 2: Confusion Matrix

A confusion matrix is a table used to visualize the performance of a classification machine learning model. It shows how many times the model correctly or incorrectly predicted each class. This helps identify areas for improvement, like high false positives for a specific class, allowing data scientists to refine the model for better accuracy. We would like to see how you can use this tool to analyze the errors of your model.

Q2.1

Get the MNIST data using `fetch_openml`. Since the data is already shuffled, take the first 30,000 elements as your train dataset.

```
12  """ Confusion Matrix
13
14  ## 1. Get the MNIST data using fetch_openml. Since the data is already shuffled, take the first 30,000 elements as your train dataset.
15
16  1. get the MNIST data.
17
18  """ **TO DO: FILL THE BLANK LINE**
19  """
20
21  from sklearn.datasets import fetch_openml
22
23  mnist = fetch_openml('mnist_784', version=1)
24
25  X, y = mnist.data, mnist.target
26  X
27
28  """since the data is already shuffled take the first 30000 elements as your train dataset.
29
30  """ **TO DO: FILL THE BLANK LINE**
31  """
32
33  X_train, X_test, y_train, y_test = X[:30000], X[30000:], y[:30000], y[30000:]
34
35  """Use the SGD classifier of Scikitlearn to classify the digits, the random state is set for the sake of reproducibility."""
36
37  from sklearn.linear_model import SGDClassifier
38  sgd_clf = SGDClassifier(random_state=42)
39  sgd_clf.fit(X_train, y_train)
```

Figure 1: Code Snippet

Q2.2

Preprocess your data using the Standard Scalar preprocessor. First, considering the nature of your data (pixels) and your classifier (SGD), explain why Standard Scalar is a good choice for data pre-processing. Next, compare the Standard Scalar with Min-Max Scaling and with no processing at all, and explain why they might not be a good choice for preprocessing.

```
41 """# 2. Use the provided code to preprocess the data and train your model.
42
43 ### **TO DO: FILL THE BLANK LINES**
44 ---
45
46 from sklearn.preprocessing import StandardScaler
47
48 scaler = StandardScaler() #----- fill this line
49 X_train_scaled = scaler.fit_transform(X_train.astype("float64"))
```

Figure 2: Code Snippet

Comment: explain why Standard Scalar is a good choice for data preprocessing

Comment: compare the Standard Scalar with Min-Max Scaling and with no processing at all, and explain why they might not be a good choice for preprocessing.

Q2.3

Make a colored diagram of the confusion matrix using ConfusionMatrixDisplay. Explain the result. Which numbers are getting misclassified? Is there any specific correlation between any 2 digits?

```
51 """# 3. Make a colored diagram of the confusion matrix using ConfusionMatrixDisplay. Explain the result. Which numbers are getting misclassified? Is there any specific correlati
52
53 from sklearn.model_selection import cross_val_predict
54
55 ##### **TO DO: FILL THE BLANK LINE*****
56
57 from sklearn.metrics import ConfusionMatrixDisplay
58
59 y_train_pred = cross_val_predict(sgd_clf, X_train_scaled, y_train, cv=3)
60 plt.rc('font', size=9)
61 ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred, display_labels=sgd_clf.classes_, cmap='viridis') #----- fill here
62 plt.show()
```

Figure 3: Code Snippet




Figure 4: Confussion Matrix

Comment: Explain the result.

Comment: Which numbers are getting misclassified? Is there any specific correlation between any 2 digits?

Q2.4

Normalize the confusion matrix by dividing each value by the total number of images in the corresponding (true) class. How does this help you to get a better understanding of the errors of your model?

```
64 """##4. Normalize the confusion matrix by dividing each value by the total number of images in the corresponding (true) class. How does this help you to get a better understandin
65
66 """ ##TO DO: FILL THE BLANK LINE""
67
68
69 plt.rc('font', size=10)
70 ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred, display_labels=sgd_clf.classes_, cmap='viridis', normalize='true') #----- fill here
71
72 plt.show()
```

Figure 5: Code Snippet




Figure 6: Normalized Confusion Matrix

Q2.5

Although it is not immediately clear from the diagram in the previous section, closer examination reveals that many digits have been incorrectly identified as 8s. To highlight these errors more, one may assign zero weight to the correct predictions. Make the errors in the confusion matrix more significant and put zero weight on the correct predictions (use the `sample_weight` feature). Plot out the updated confusion matrix and explain what you observe.

```

74 """# 5. Make the errors more significant and put zero weight on the correct predictions. Explain your results.
75
76 *** **TO DO: FILL THE BLANK LINE**
77 ***
78
79 sample_weight = (y_train_pred != y_train)
80 plt.rc('font', size=10)
81 ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred,
82                                       display_labels=sgd_clf.classes_,
83                                       normalize="true", values_format=".05")
84 plt.show()

```

Figure 7: Code Snippet




Figure 8: Corrected Confussion Matrix

Conclusion

In conclusion, the MNIST dataset has been successfully loaded from the source on internet. After certain preprocesses to the data, a linear classifier with SGD (Stochastic Gradient Descent) has been used for training a model. A part of the data tested the model. The results constituted a confusion matrix and the matrix has been analyzed with different versions.

Question 3: Nearest Neighbors from Scratch

In this exercise, you will complete the implementation of the k-Nearest Neighbours (KNN) algorithm in the skeleton code. Do not change the regression, classification dataset and query values for the final submission.

Q3.1

Implement the function to calculate the Euclidean distance between two points.

```
46 def euclidean_distance(point1, point2):  
47     # implement this function to return euclidean distance between point1 and point2  
48     return abs(point1-point2)
```

Figure 1: Code Snippet for Euclidean Distance Function

Q3.2

Implement the functions to calculate mean and mode which will serve as choice functions for the KNN model. For regression tasks, use the mean; for classification tasks, use the mode.

```
50 def mean(labels):  
51     # implement this function to return the mean of the labels.  
52     sum = 0  
53     for val in labels:  
54         sum += val  
55  
56     return sum/len(labels)
```

Figure 2: Code Snippet for Mean Function

```
58 def mode(labels):  
59     # implement this function to return the mode of the labels.  
60     occurrences = {}  
61     for val in labels:  
62         if(val in occurrences.keys()):  
63             occurrences[val] += 1  
64         else:  
65             occurrences[val] = 1  
66  
67     max_val = 0  
68     max_occ = 0  
69     for val in occurrences.keys():  
70         if(occurrences[val] > max_occ):  
71             max_val = val  
72             max_occ = occurrences[val]  
73  
74     return max_val
```

Figure 3: Code Snippet for Mode Function

Q3.3

Complete the knn function by following the comments provided in the skeleton code and run your custom KNN on your regression and classification data to find the height of the person who is 55 years old and whether an 18-year-old likes pineapple or not.

```
12 def knn(data, query, k, distance_fn, choice_fn):
13     """
14     """
15     neighbor_distances_and_indices = [] # in order with the dataset
16
17     # Calculate the distance between the query example and all the examples in the data.
18     index=0
19     for val in data.iloc[:,0]:
20         neighbor_distances_and_indices.append(list([euclidean_distance(val, query), index]))
21         index += 1
22
23     # Sort the distances and return the labels of the k nearest neighbors.
24     neighbor_distances_and_indices.sort()
25
26     # Pick the first k entries from the sorted collection
27     k_nearest_neighbors = neighbor_distances_and_indices[:k]
28
29     # Get the labels of the selected k entries
30     k_nearest_labels = []
31     for ls in k_nearest_neighbors:
32         k_nearest_labels.append(data.iloc[ls[1], 1])
33
34     # If regression (mean), if classification (mode)
35     return choice_fn(k_nearest_labels)
```

Figure 4: Code Snippet for KNN Function

```
121 ## Load the Regression Data. The first index consists of age(feature) and the second index is the label. The label is height of the person in
122 import pandas as pd
123 import numpy as np
124
125 reg_data = pd.read_csv('ece657_a2_q3_data/regression_data.csv') # Load the data
126
127 reg_query = np.array([[55]]) # reshape to fit scikit-learn requirements
128
129 # Custom KNN Prediction
130 custom_reg_prediction = knn(reg_data, reg_query, k=3, distance_fn=euclidean_distance, choice_fn=mean)
131
132 # print(f"The predicted height of {reg_query[0][0]} years old person: {custom_reg_prediction} cm")
133 print("Custom KNN Regression Prediction:", custom_reg_prediction)
```

Figure 5: Code Snippet for Regression Prediction

```
141 # Load the Classification Data. The first index consists of age(feature) and the second index is the label. The label 0 is for likes pineapple
142 clf_data = pd.read_csv('ece657_a2_q3_data/classification_data.csv') # Load the data
143
144 clf_query = np.array([[18]]) # reshape to fit scikit-learn requirements
145
146 # Custom KNN Prediction
147 custom_clf_prediction = knn(clf_data, clf_query, k=3, distance_fn=euclidean_distance, choice_fn=mode)
148
149 # print(f"Would a {clf_query[0][0]} years old person like pineapples: {bool(custom_clf_prediction)}")
150 print("Custom KNN Classification Prediction:", custom_clf_prediction)
```

Figure 6: Code Snippet for Classification Prediction

Q3.4 (NOT FINISHED)

Using KNN library from scikit-learn. In the skeleton code, add code to run the same datasets with the same k value for KNeighborsRegressor and KNeighborsClassifier. Did you get the same value for regression and classification for the datasets provided in the skeleton code?

```
75 def sklearn_knn_regression(reg_data, reg_query):
76     # Initialize the KNN regressor with 3 nearest neighbors
77     knn_reg = KNeighborsRegressor(n_neighbors=3)
78
79     # Fit the model on the training data; use all but the last column as features and the last column as the target
80     knn_reg.fit(reg_data.iloc[:, :-1], reg_data.iloc[:, -1])
81
82     # Predict the output for the provided query and return the first (and likely only) prediction
83     skl_reg_prediction = knn_reg.predict(reg_query)
84
85     return skl_reg_prediction
```

Figure 7: Code Snippet for Scikit-Learn Regression Function

```
139 # Scikit-learn KNN Regression
140 skl_reg_prediction = sklearn_knn_regression(reg_data, reg_query)
141
142 print("Scikit-learn KNN Regression Prediction:", skl_reg_prediction)
```

Figure 8: Code Snippet for Scikit-Learn Regression Prediction

```
87 def sklearn_knn_classification(clf_data, clf_query):
88     # Initialize the KNN classifier with 3 nearest neighbors
89     knn_clf = KNeighborsClassifier(n_neighbors=3)
90
91     # Fit the model on the training data; use all but the last column as features and the last column as the target
92     knn_clf.fit(clf_data.iloc[:, :-1], clf_data.iloc[:, -1])
93
94     # Predict the class for the provided query and return the first (and likely only) prediction
95     skl_clf_prediction = knn_clf.predict(clf_query)
96
97     return skl_clf_prediction
```

Figure 9: Code Snippet for Scikit-Learn Classification Function

```
155 # Scikit-learn KNN Classification
156 skl_clf_prediction = sklearn_knn_classification(clf_data, clf_query)
157
158 print("Scikit-learn KNN Classification Prediction:", skl_clf_prediction)
```

Figure 10: Code Snippet for Scikit-Learn Classification Prediction

- Custom KNN Regression Prediction: 128.24666666666667
- Scikit-learn KNN Regression Prediction: [128.24666667]

- Custom KNN Classification Prediction: 0
- Scikit-learn KNN Classification Prediction: [0]
- Custom KNN Classification Prediction: 0
- Scikit-learn KNN Classification Prediction: [0]
- Prediction for weighted KNN: 1

Q3.5

Complete the `weighted_mode` and `knn_weighted` functions in the provided skeleton code to classify whether a 15-year-old likes pineapple or not. For the `KNeighborsClassifier` assign weights proportional to the inverse of the distance from the query point and then classify.

```
102 from collections import defaultdict
103
104 def weighted_mode(labels, weights):
105     # Initialize a defaultdict to store the sum of weights for each label
106     sum_weighted_labels = defaultdict(int)
107
108     # Iterate through each label in the labels list and Sum the weights for each label and store in the defaultdict
109     for l in labels:
110         sum_weighted_labels[l] += weights[l]
111
112     # Determine the label with the maximum sum of weights
113     max_weighted_label = max(sum_weighted_labels, key=sum_weighted_labels.get)
114
115     # Return the label that has the highest sum of weights
116     return max_weighted_label
```

Figure 11: Code Snippet for `weighted_mode` Function

```
118 def knn_weighted(data, query, k, distance_fn, choice_fn, weights):
119     neighbor_distances_and_indices = []
120
121     # Calculate the distance between the query example and all the examples in the data.
122     index=0
123     for val in data.iloc[:,0]:
124         neighbor_distances_and_indices.append(list([euclidean_distance(val, query), index]))
125         index += 1
126
127     # Sort the distances and return the labels of the k nearest neighbors.
128     neighbor_distances_and_indices.sort()
129
130     # Pick the first k entries from the sorted collection
131     k_nearest_neighbors = neighbor_distances_and_indices[:k]
132
133     # Get the labels of the selected k entries
134     k_nearest_labels = []
135     for ls in k_nearest_neighbors:
136         k_nearest_labels.append(data.iloc[ls[1], 1])
137
138     # Apply the weighted mode function and return nearest neighbors too
139     return choice_fn(k_nearest_labels, weights)
```

Figure 12: Code Snippet for `knn_weighted` Function

```
177 weights = {0: 1, 1: 2}
178
179 # Query for whether a 15-year-old likes pineapple or not. The classification should be 1 as this exact sample is present in the dataset
180 # but because of unbalanced dataset this will be predicted as class 0.
181
182 clf_query = np.array([[15]]) # reshape to fit scikit-learn requirements
183
184 # Custom KNN Prediction
185 custom_clf_prediction = knn(clf_data, clf_query, k=3, distance_fn=euclidean_distance, choice_fn=mode)
186
187 # Scikit-learn KNN Classification
188 skl_clf_prediction = sklearn_knn_classification(clf_data, clf_query)
189
190 print("Custom KNN Classification Prediction:", custom_clf_prediction)
191 print("Scikit-learn KNN Classification Prediction:", skl_clf_prediction)
192
193 clf_prediction = knn_weighted(clf_data, clf_query, k=3, distance_fn=euclidean_distance, choice_fn=weighted_mode, weights=weights)
194 print("Prediction for weighted KNN:", clf_prediction)
```

Figure 13: Code Snippet for Weighted Classification Prediction

Q3.6

In Scikit-learn, several algorithms involve randomness in their operation and therefore provide a `random_state` parameter to ensure reproducibility of results by controlling the seed of the random number generator. Do `KNeighborsRegressor` and `KNeighborsClassifier` have it?

Yes, both `KNeighborsRegressor` and `KNeighborsClassifier` in scikit-learn do have a parameter for controlling the seed of the random number generator to ensure reproducibility of results. This parameter is typically named `random_state`. Example: `regressor = KNeighborsRegressor(n_neighbors=3, random_state=42)`

Question 5: Data Visualization and Model Selection

Consider the dataset found here. It is a modified version of the popular Heart Disease Dataset originally comprising 303 instances with 13 features, and serves as a multivariate resource for classification tasks in the realm of health and medicine, featuring categorical, integer, and real feature types.

Q5.1

Display basic information about the dataset, including the data types of each column followed by a summary of the dataset's statistics, including count, mean, standard deviation and max values for each column.

```
C:\Users\mapel\OneDrive - University of Waterloo\ECE657\A2>py ece657_a2_q5.py
Basic Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 549 entries, 0 to 548
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0    cp          549 non-null    int64
1    ca          549 non-null    int64
2    thalach     549 non-null    int64
3    oldpeak     549 non-null    float64
4    thal        549 non-null    int64
5    target      549 non-null    int64
dtypes: float64(1), int64(5)
memory usage: 25.9 KB
None

First few rows of the dataset:
   cp  ca  thalach  oldpeak  thal  target
0   0   3     145      6.2     3        0
1   0   1     141      2.8     3        0
2   0   1     156      0.1     3        0
3   0   2      90      1.0     1        0
4   0   2     165      1.0     3        0

Summary Statistics:

```

	cp	ca	thalach	oldpeak	thal	target
count	549.000000	549.000000	549.000000	549.000000	549.000000	549.000000
mean	0.561020	1.078324	140.845173	1.492532	2.504554	0.094718
std	0.942901	1.039048	23.077379	1.288169	0.685138	0.293092
min	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	125.000000	0.300000	2.000000	0.000000
50%	0.000000	1.000000	143.000000	1.200000	3.000000	0.000000
75%	1.000000	2.000000	160.000000	2.400000	3.000000	0.000000
max	3.000000	4.000000	195.000000	6.200000	3.000000	1.000000

Figure 1: Basic information about the dataset

Q5.2

Calculate and display the proportion of each unique value in the 'target' column. This analysis will help you understand the frequency of each category within the target variable, providing insights into the dataset's balance. Is the dataset imbalanced?

```
Distribution of target variable:
target
0    497
1     52
Name: count, dtype: int64

Proportion of each unique value in the target variable:
target
0    0.905282
1    0.094718
Name: proportion, dtype: float64
```

Figure 2: Proportion of each unique value in the target

Because the difference in proportions of the target values is too big, we can call the dataset is imbalanced.

Q5.3

Plot histograms for all numerical attributes in the dataset. You will identify all numerical columns and generate histograms with specified bin sizes to visually assess the distribution of these attributes. This visualization is vital for spotting any skewness or outliers in the data and understanding the distribution of numerical variables. Which of the features are skewed? Also, mention if they are left-skewed or right-skewed.

```
Skewness of numerical features:
cp          1.369990
ca          0.587301
thalach    -0.327264
oldpeak     0.790790
thal       -1.212341
target      2.775679
dtype: float64
```

Figure 3: Skewness of numerical values

From the values for skewness, we can say that features cp, ca, oldpeak,

and `target` are all right-skewed. `thal` is a left-skewed feature whereas `thalach` is appricimately symmetric.

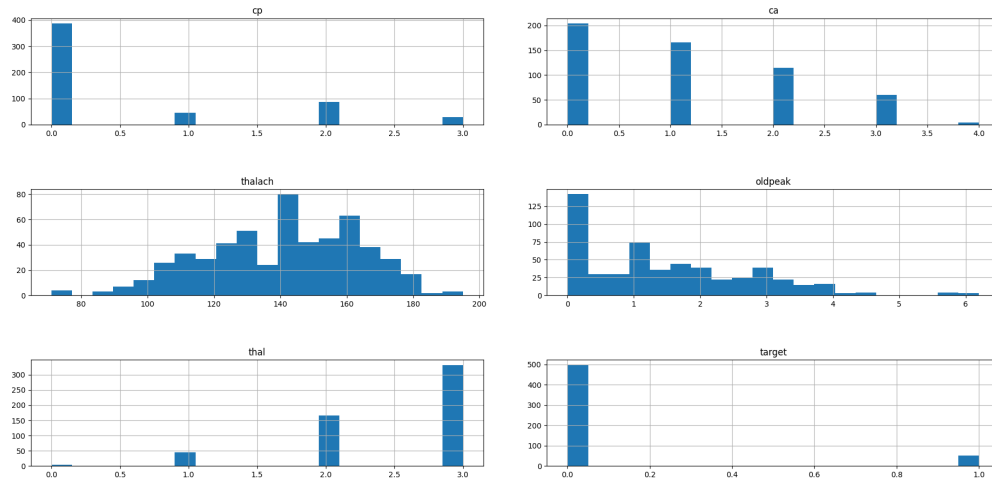


Figure 4: Histograms for attributes in the dataset

Q5.4

Generate a heatmap to visualize the correlations between different variables. Understanding these correlations is crucial for identifying relationships between variables, which can inform feature selection and predictive modelling strategies. Is there a high correlation with features or with the target variable?



Figure 5: Correlation Matrix Heatmap

When we check the correlation values on the matrix, it can be concluded that the correlations between features and the target variable are relatively weak.

Q5.5

Scale the features of the dataset using Python and pandas along with scikit-learn's StandardScaler. Remove the target column and apply scaling to the remaining features to standardize them, converting the scaled data back into a DataFrame with the original column names. Why is this step important? Report the mean and standard deviation of the scaled features.

Scaling is important to ensure that each feature contributes to the model equally. Otherwise, some features may end up dominate the others.

```
Mean of scaled features:
cp      2.912060e-17
ca      -2.264936e-17
thalach -1.132468e-16
oldpeak  9.059743e-17
thal     -6.471245e-17
dtype: float64

Standard deviation of scaled features:
cp      1.000912
ca      1.000912
thalach 1.000912
oldpeak 1.000912
thal     1.000912
dtype: float64
```

Figure 6: The mean and standard deviation of the scaled features

Q5.6

Assign the scaled features to X and the target variable to y . Specify the size of the test set to be 20% of the entire dataset, ensure the data is split in a way that maintains the proportion of classes in both training and testing sets by stratifying on y , and set a random state to 25 for reproducibility. Print the shapes of the training and testing sets.

```
Shapes of training and testing sets:
X_train: (439, 5)
X_test: (110, 5)
y_train: (439,)
y_test: (110,)
```

Figure 7: The shapes of training and testing sets

Q5.7

Table 2: Classifier Parameters for Different Sets

Classifier	Set 1	Set 2	Set 3
KNN	n_neighbors=3, weights='uniform'	n_neighbors=10, weights='uniform', p=2	n_neighbors=15, weights='distance', p=1
Decision Tree	max_depth=3, min_samples_split=2, random_state: 25	max_depth=10, min_samples_leaf=2, random_state: 25	criterion='entropy', max_depth=20, min_samples_leaf=4, random_state: 25
Random Forest	n_estimators=50, max_features='auto', random_state: 25	n_estimators=200, max_depth=10, max_features='log2', random_state: 25	n_estimators=350, max_depth=20, min_samples_split=3, random_state: 25
XGBoost	learning_rate=0.01, max_depth=3, n_estimators=50, random_state: 25	learning_rate=0.15, max_depth=10, n_estimators=150, random_state: 25	learning_rate=0.25, max_depth=15, n_estimators=250, random_state: 25

Figure 8: The table provided in the assignment description document

Evaluate various models with different hyperparameters listed in Table 2. Implement these models using sklearn for each set of parameters. Based on training and testing accuracy scores and confusion matrices, determine the best-performing model for health data. Additionally, provide insights into the performance of each model.

<p>Evaluating Set 1:</p> <p>KNN: Training Accuracy: 0.9749430523917996 Testing Accuracy: 0.9454545454545454 Confusion Matrix: [[100 0] [6 4]]</p> <p>Decision Tree: Training Accuracy: 0.9339407744874715 Testing Accuracy: 0.9090909090909091 Confusion Matrix: [[96 4] [6 4]]</p> <p>Random Forest: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9545454545454546 Confusion Matrix: [[100 0] [5 5]]</p> <p>XGBoost: Training Accuracy: 0.9043280182232346 Testing Accuracy: 0.9090909090909091 Confusion Matrix: [[100 0] [10 0]]</p>	<p>Evaluating Set 2:</p> <p>KNN: Training Accuracy: 0.9339407744874715 Testing Accuracy: 0.9454545454545454 Confusion Matrix: [[100 0] [6 4]]</p> <p>Decision Tree: Training Accuracy: 0.9817767653758542 Testing Accuracy: 0.9454545454545454 Confusion Matrix: [[100 0] [6 4]]</p> <p>Random Forest: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9545454545454546 Confusion Matrix: [[100 0] [5 5]]</p> <p>XGBoost: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9636363636363636 Confusion Matrix: [[100 0] [4 6]]</p>	<p>Evaluating Set 3:</p> <p>KNN: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9454545454545454 Confusion Matrix: [[100 0] [6 4]]</p> <p>Decision Tree: Training Accuracy: 0.958997722095672 Testing Accuracy: 0.9181818181818182 Confusion Matrix: [[96 4] [5 5]]</p> <p>Random Forest: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9545454545454546 Confusion Matrix: [[100 0] [5 5]]</p> <p>XGBoost: Training Accuracy: 0.9954441913439636 Testing Accuracy: 0.9636363636363636 Confusion Matrix: [[100 0] [4 6]]</p>
--	--	--

Figure 9: Results of the model evaluation

Set 2 - XGBoost stands out as the best-performing model (Even Set 3 - XGBoost has the same test accuracy, Set 2 has less max-depth):

- Testing Accuracy: 96.36
- Confusion Matrix: $\begin{bmatrix} 100 & 0 \\ 4 & 6 \end{bmatrix}$

It has a relatively balanced confusion matrix indicating that most of the test samples were correctly classified and handled class imbalances effectively; therefore, this model achieves the highest testing accuracy.

Other models:

- kNNs showed a good consistent performance with high training and testing accuracy for all sets.
- Decision trees had improvement in accuracy with deeper trees but can be prone to overfitting (in Set 2, the high accuracy of training).
- Random Forests seem to be more reliable when it comes the similarity of training and testing accuracies.

Q5.8

What are the potential challenges that exist with working with real-world healthcare data for machine learning applications?

Using real healthcare data for machine learning is challenging. One may need to deal with issues like data quality problems, integrating complex data, and handling imbalanced classes. There are also ethical concerns about privacy and bias, and the need for clear explanations makes it more complex. Technical challenges such as figuring out which features to use, making sure models are accurate, and meeting regulations make it even harder. Solving these problems requires teamwork and different areas of expertise to build trustworthy and fair healthcare solutions.